PMM Decision Intelligence Agent

Code Name: Prism · Design Specification v0.1 · Project Proposal

PMM AI agent design incorporating signal quality framework, decision context model, outcome tracking, and organizational adoption requirements.

Version	0.1
Status	Draft for Internal Review / Project Proposal
Audience	PMM Leadership · Project Sponsors
Purpose	Full design scope

Executive Summary

The PMM Decision Intelligence Agent (DIA) is a signal-processing and decision-forcing system designed to eliminate the gap between market signal emergence and PMM strategic action. It is not a content generator, insight dashboard, or analytics platform. Its sole purpose is to force decisions earlier, long before competitors define the narrative, before field misalignment compounds, and before a positioning shift becomes a reactive catch-up exercise.

Problem Statement

Product Marketing operates downstream of market reality. Competitive shifts are observed but not acted on decisively. Field signals are fragmented across systems, ISC notes, EPM data, MCC streams, sales call recordings, Slack channels, analyst papers, competitor content, media channels, BP feedback, etc., and by the time these signals converge into an actionable view, the market has already moved.

The consequence is not a lack of information. PMM has access to more signals than any previous era. The consequence is decision latency: the gap between when a shift becomes visible in the data and when PMM commits to a repositioning response. That gap is where competitors gain ground and where field teams develop misaligned narratives in the absence of clear PMM direction.

Root Causes

- Signals are fragmented across incompatible systems with no unified interpretation layer
- Pattern detection is manual, inconsistent, and dependent on individual analyst attention
- Narrative misalignment is detected late, often after it has embedded in field behavior
- There is no mechanism that distinguishes "we've seen this signal and decided to hold" from "we haven't processed this signal yet"
- Decision deferral is costless in the short term and invisible until the window has closed

The Specific Failure Mode DIA Addresses

DIA is not designed to solve the information problem. It is designed to solve the decision-timing problem. The system's value is not in surfacing signals PMM couldn't find; it's in forcing a response to signals PMM has already seen but hasn't acted on. The difference between a functioning DIA and a well-maintained PowerPoint, Excel spreadsheet, or dashboard is that DIA assigns a commit-by date, names an owner, and logs whether the decision was made.

Design Philosophy

The Decision Forcing Function

Every positioning decision ultimately reduces to a choice between two directions. Sequencing, emphasis, and timing are real dimensions of that choice, but they do not change the fact that at some point, a commitment must be made. The system is designed to force that commitment before it becomes irreversible by default.

Decision Types

Not all decisions require the same magnitude of response. Treating a framing adjustment as a full platform repositioning wastes execution capacity and creates organizational noise. The system classifies each decision by type before forcing a choice:

Decision Type	Definition	Implication
FULL REPOSITION	Core platform narrative changes. Affects all external materials, sales motion, marketing campaign programs, field marketing tactics, AR/PR strategies, etc.	High cost, high consequence. Requires senior sign-off.
MESSAGING ACCENT	Shift emphasis within existing narrative. Proof points updated; core claim intact.	Moderate effort. Can execute within a campaign cycle.
FRAMING SHIFT	Same proof points, reframed for a different audience or competitive context.	Low effort. Talk track and battlecard update.
HOLD WITH TRIGGER	Maintain current approach. A defined signal or event will force re-evaluation.	Active choice. Not passive inaction.

Design Rule: Hold is Not Neutral
A decision to hold current positioning must specify what would change that decision. If no trigger is defined, the system does not accept B as a valid response. An unresolved trigger after the commit-by date is logged as DEFERRED (a distinct status that is visible to PMM leadership).

Architecture Overview

The system processes signals through six layers, each with a distinct function. The addition of a Signal Quality Framework (Layer 2), between ingestion and embedding, prevents pattern contamination from low-independence or adversarial signals.

```
Signals
   │
   ▼
Layer 1: Signal Ingestion  ──────────── Competitor content, analyst research,
   │                                    press feeds, Salesforce, Slack, call
   │                                    transcripts, field plans, win/loss, etc.
   ▼
Layer 2: Signal Quality Framework ───── Independence scoring, recency weighting,
   │                                    foe discount, Signal Tier assignment
   │
   ▼
Layer 3: Embedding + Clustering ──────── SBERT → UMAP → HDBSCAN → BERTopic
   │
   ▼
Layer 4: Interpretation Engine ───────── Cluster → Pattern statement (X → Y)
   │
   ▼
Layer 5: Decision Engine (Core) ──────── Pattern + Constraint config → Decision
   │                                    Output with type, timing, owner, trigger
   ▼
Layer 6: Delivery ─────────────────────── Slack (Prism), weekly digest,
                                          Confluence knowledge base
```

System Architecture — Layer Detail

Layer 1: Signal Ingestion

External Signals

- Competitor web content (scraped, structured CSS selectors, not full-page dumps)
- Analyst research (PDF ingestion via automated RSS and manual upload)
- Press releases and trade media (RSS feeds, curated source list)
- Competitor pricing and packaging changes (web monitoring)

Internal Signals

- Salesforce ISC: opportunities, deal notes, stage progression
- Sales call transcripts (where access is available)
- Slack/Teams: selected channels and distribution groups (export, not live)
- Field marketing plans, programs, and reports
- Win/loss reports

> **V1 Scope Discipline**
> Not all signal sources will be accessible in V1. Phase 0 (Week 1–2) exists specifically to map what is actually accessible before any pipeline is built. Sources requiring >4 weeks of IT or legal access negotiation are V2 by default.

Layer 2: Signal Quality Framework

This layer exists to prevent the interpretation and decision layers from processing contaminated signal.

Source Independence Scoring

Source Category	Independence Weight	Rationale
Analyst reports / research	**1.0x**	Independent, methodologically grounded. Lagging but credible.
Sales call transcripts	**1.0x**	Unfiltered buyer signal. High value.
Field team feedback / win-loss	**1.0x**	Outcome-validated, direct market contact.
Internal Slack / channels	**0.75x**	Reactive to external signals. Useful for confirmation, not discovery.
Press / RSS feeds	**0.5x**	Often reflects competitor PR. Corroborating only.
Competitor website content	**0.25x**	Adversarial by design. Discounted, not excluded.

Multiple posts from same source	**Counts as 1 signal regardless of volume**	Same-origin does not create independence.

Signal Tier System

Signal Tiers are auditable, explainable, and do not imply false precision. A tier is assigned based on observable criteria, not model self-assessment.

Tier	Signals (weighted)	Source Categories	Recency	Action
EMERGING	≥2	1+	<14 days	Monitor. Do not trigger decision.
CONFIRMED	≥3	2+	<21 days	Trigger decision model.
ESTABLISHED	≥5	3+	30+ days sustained	Escalate urgency to HIGH.

Recency Weighting

Signal Age	Weight	Status
0–7 days	**1.0x**	Full weight
8–21 days	**0.75x**	Decaying
22–60 days	**0.5x**	Corroborating only
>60 days	**0.0x**	Archive, does not contribute to tier

Layer 3: Signal Embedding + Clustering

Unstructured signal text is converted to dense vector representations, reduced in dimensionality, and clustered to identify co-occurring semantic patterns. The pipeline is:

```
Raw Text Documents
     │
     ▼  SBERT (sentence-transformers)
Dense Embeddings
     │
     ▼  UMAP (dimensionality reduction)
Reduced Vector Space
     │
     ▼  HDBSCAN (density-based clustering)
Signal Clusters
     │
     ▼  c-TF-IDF / BERTopic
Topic Labels + Directional Summaries
```

> **Manual Review Gate**
> For the first 30 days of V1 operation, all cluster outputs are reviewed manually before triggering decision logic. This calibration period establishes baseline pattern quality and prevents spurious triggers from contaminating the decision record.

Layer 4: Interpretation Engine

Clusters are passed to an LLM with a constrained prompt that forces a single directional pattern statement. The output format is non-negotiable: "Market is shifting from [X] to [Y] because [reason]." No lists, no summaries, no options.

Prompt 1: Pattern Compression

```
You are identifying structural shifts in market behavior.
Input:  [Cluster of signals]
Task:   Reduce these signals into ONE pattern describing a directional shift.
Rules:  Do not summarize. Do not list. Do not hedge.
Output: "Market is shifting from [X] to [Y] because [reason]."
Reject: Any output that does not follow this exact format.
```

Layer 5: Decision Engine

The decision engine applies constraint logic against the current positioning assumption configuration. When a pattern from Layer 4 overlaps with an encoded assumption, the engine generates a Decision Trigger output. The output format is described in the Decision Output Model section.

Positioning Assumption Config: Governance

The positioning assumption config is the most fragile component of the system. It is manually defined and must remain current to produce valid decisions. In the first release, this config is treated as a governed artifact:

- Version-controlled in an appropriate repository with a required change log entry for every edit
- Each assumption carries: id, statement, date_encoded, last_reviewed, owner
- Assumptions not reviewed in >90 days are flagged STALE and excluded from decision logic until reviewed
- Quarterly config review is **mandatory** and is a V1 requirement, not V2

Prompt 2 : Constraint Identification

```
Input:  [Pattern statement] + [Current positioning assumption]
Task:   Identify which assumption no longer holds.
Output: "This model assumes [X]. This no longer holds because [Y]."
```

Prompt 3: Decision Enforcement

```
  Input:  [Constraint output]
  Task:   Force a decision. No option lists.
  Output: "This requires a decision: [A] or [B]."
  Note:   A and B are concrete, named actions, not abstract directions.
```

Prompt 4: Impact Framing

```
  Input:  [Decision]
  Task:   Identify where current narrative or strategy breaks if no decision is
made.
  Output: "If unchanged, this will impact [area] within [timeframe]."
```

Layer 6: Delivery

- “Prism” Slack bot: real-time decision triggers pushed to named PMM channels
- “Prism” weekly email digest: summary of open, decided, and deferred decisions
- “Prism” knowledge base: Confluence archive of all decision triggers and outcomes

Decision Output Model

The Decision Output Model is the core interface between the system and the PMM team. Every output from Layer 5 uses this exact format. The format is not a template; it is a contract. Deviations indicate system failure, not flexibility.

Output Format Specification

```
━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━
  DECISION TRIGGER: [Model Name] — [Subject]
  STATUS: ACTION REQUIRED  |  MONITORING  |  RESOLVED
━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━

  SIGNAL BASIS
    Signals detected:    [N] (weighted) over [X] days
    Source diversity:    [N] source categories
    Independence score:  HIGH | MEDIUM | LOW
    Signal tier:         EMERGING | CONFIRMED | ESTABLISHED

  CONTEXT
    Decision type:       [FULL REPOSITION | MESSAGING ACCENT | FRAMING SHIFT
                          | HOLD WITH TRIGGER]
    Timing window:       [Specific constraint (e.g. before campaign brief locks)]
    Commit by:           [Auto-calculated date]
    Decision owner:      [Named role, not "PMM"]

  THE DECISION
    A)   [Option A: specific, concrete action]

    B)   [Option B: specific, concrete action OR hold]
         REQUIRED if B is HOLD: Define re-evaluation trigger
         ↳ Re-evaluate if: [specific signal or event]
            within [timeframe], OR decision re-opens automatically.

  IMPACT (if no decision made by commit-by date)
    Area:        [narrative | product | pricing | field alignment]
    Timeframe:   [short | medium | long]
    Consequence: [one sentence: what breaks, for whom, by when]

  OUTCOME LOG
    Status:        [ OPEN | DECIDED-A | DECIDED-B | DEFERRED | EXPIRED ]
    Date decided:  ___________
    Notes:         ___________
    30-day review: [ ] Signal accurate  [ ] Inaccurate  [ ] Inconclusive
━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━
```

Worked Example: IBM Fusion AI Positioning

```
━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━
  DECISION TRIGGER: Competitive Axis Shift — Fusion AI Positioning
  STATUS: ACTION REQUIRED
━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━
```

```
SIGNAL BASIS
  Signals detected:    4 (weighted) over 14 days
  Source diversity:    3 categories (analyst, field, competitor)
  Independence score:  HIGH — cross-category confirmation
  Signal tier:         CONFIRMED

CONTEXT
  Decision type:       MESSAGING ACCENT SHIFT
                       (Reframe existing proof points; core narrative intact)
  Timing window:       Must decide before Q3 campaign brief locks — June 14
  Commit by:           June 7  (7-day execution lead time)
  Decision owner:      PMM Lead, Platform Narrative

THE DECISION
  A)  Shift Fusion messaging to emphasize operational AI readiness.
      Maintain platform narrative; update proof points and talk tracks
      to reflect production-stage AI workloads.

  B)  Hold current exploratory AI framing.
      ↳ Re-evaluate if: >2 field calls flag "too early-stage" objection
         within 30 days, OR any Tier 1 analyst publishes "operational AI
         maturity" as a primary buying criterion.

IMPACT (if no decision by June 7)
  Area:        AI, platform, and sovereignty messaging
  Timeframe:   Medium (Q3 exposure)
  Consequence: Fusion positioned as exploratory in a market that has moved
               to operational accountability framing visible to enterprise
               architects before Q3 campaigns launch.

OUTCOME LOG
  Status:        [ OPEN ]
  Date decided:  ___________
  Notes:         ___________
  30-day review: [ ] Accurate  [ ] Inaccurate  [ ] Inconclusive
```

Decision Models V1

Three decision models are scoped for V1. Each model has a defined purpose, input sources, detection logic, and a threshold before a decision trigger is emitted. All three are encoded with the Decision Output format and governed by the positioning assumption config.

Decision Model 1: Narrative Drift

Purpose: Detect when Fusion is being repositioned by the field as a product (storage, virtualization) rather than a platform. This is an internal signal model — it detects PMM message decay in execution, not a competitive shift.

Inputs	ISC call notes, sales call transcripts, Slack seller channels, field content
Detection	IF >30% of [product] references = [use-case list] AND [product focus terms] decline over same period
Threshold	Signal tier reaches CONFIRMED (3+ weighted signals, 2+ source categories, <21 days)
Output	Signal: Field positioning converging on product-led messaging; Constraint: [product] must be positioned as a [this], not [that]; Decision: Reinforce [product] narrative with field enablement OR accept [new positioning]; If hold: Re-evaluate if >3 additional field calls reflect product framing within 30 days

Decision Model 2: Entry Point Validity

Purpose: Identify whether the job-to-be-done opening deals matches the entry point PMM is emphasizing. This model surfaces misalignment between marketing messaging and sales reality.

Inputs	ISC opportunity create, deal notes, win/loss summaries
Detection	Cluster deal drivers: [list (examples - VMware Transition, AI Infrastructure, Storage Modernization, Hybrid Complexity)]; IF >50% of new pipeline driven by ONE category AND PMM messaging emphasizes a different entry point
Threshold	Minimum 20 deals analyzed; signal tier CONFIRMED
Output	Signal: [Category] is driving majority of pipeline entry Constraint: Current marketing emphasizes [different entry point] Decision: Pivot entry strategy to [dominant category] OR maintain current positioning If hold: Re-evaluate at next quarterly pipeline review or if pipeline share shifts by >15%

Decision Model 3: Competitive Axis Shift

Purpose: Detect when the basis of competition changes, not just which competitors are winning, but what they are competing on. An axis shift means the criteria buyers use to evaluate a product has changed, which requires a narrative response regardless of feature parity.

Inputs	Competitor websites, analyst reports, curated media outlets, pricing and packaging signals
Detection	Track directional axes (example: cost vs. control, flexibility vs. governance, infrastructure vs. AI system readiness, etc.); IF ≥3 weighted independent competitors shift messaging direction AND analyst language confirms the directional trend
Threshold	Signal tier CONFIRMED; adversarial discount applied to all competitor signals
Output	Signal: Market shifting from [X axis] to [Y axis] Constraint: Fusion narrative positioned on [X axis framing] Decision: Reposition toward [Y axis] OR defend [X axis] as differentiated position If hold: Define the analyst trigger or deal loss pattern that would force re-evaluation

Outcome Tracking

Outcome tracking is a V1 requirement, not a V2 enhancement. Without a feedback loop from day one, the system cannot be calibrated, and there is no evidence base for proving or disproving its value. A system that produces decisions with no record of what was decided or what happened is, in practice, a dashboard.

Decision Status

Status	Definition
OPEN	Decision trigger has been delivered. Awaiting response from decision owner.
DECIDED-A	Decision owner chose Option A. Action is underway.
DECIDED-B	Decision owner chose Option B (hold with trigger). Monitoring trigger is active.
DEFERRED	Commit-by date passed with no response logged. Visible to leadership. Not neutral.
EXPIRED	Timing window closed. Decision is no longer actionable. Post-mortem logged.

Review Cadence

For each closed decision (DECIDED-A, DECIDED-B, or EXPIRED), a structured review is conducted at 30, 60, and 90 days:

- 30 days: Was the signal accurate? Did the pattern described by Layer 4 hold?
- 60 days: Was the decision response (A or B) appropriate given what is now known?
- 90 days: Did the predicted impact materialize? What was the actual consequence?

Review outcomes are classified as Accurate, Inaccurate, or Inconclusive and feed back into threshold calibration. A model with >40% Inaccurate reviews in a rolling 90-day window is suspended pending review of its detection logic and positioning assumption config.

Organizational Adoption

This is the section most likely to determine whether this product produces value. The technical pipeline can work perfectly and still fail if the organizational structure around it is not designed to force response. A system that produces decisions that no one is accountable for is no system at all.

> **The Critical Question**
> Before any code is written, PMM leadership must answer one question honestly: will the PMM team act on forced binary decisions from a system, or will they rationalize around them? The answer to that question determines whether this is worth building. This is an organizational question, not a technical one.

Requirements Before Build Begins

1. PMM leadership formally designates a decision owner for each of the three V1 decision models. These are named individuals or specific roles, not "the PMM team."
2. A decision SLA is agreed: 48-hour acknowledgment, 5 business days to log a status (DECIDED-A, DECIDED-B, or a documented dispute).
3. Leadership adopts the "no action is a choice" norm: a DEFERRED status is not neutral. It is visible in the weekly digest and in leadership reporting.
4. A governance compact is signed: “Prism” decision triggers require a logged response. Teams may disagree and choose B or document a rebuttal, but silence is not an acceptable response.

Without these four commitments in place before Phase 1 begins, the build should not proceed. The system's value is entirely contingent on behavioral response. A technically complete system deployed into an organization that treats its outputs as optional produces no value and significant sunk cost.

Quarterly Calibration Review

Every quarter, PMM leadership reviews the decision record as a group:

- Which decisions were acted on? What was the outcome?
- Which decisions were overridden? Why? Was that call right?
- Which models are producing accurate signals vs. noise?
- Which positioning assumptions in the config need to be updated?

This review is not a system audit; it is a strategic forcing function in its own right. It makes the cost of inaction visible and creates accountability for the quality of the decisions the team is and is not making.

Technology Stack

The architecture is designed to be model-agnostic at the LLM layer. IBM watsonx.ai is the current internal preference, but the system's prompt structure and output format constraints are identical regardless of which model is used. If pattern quality is insufficient with the preferred model, a swap should not require architectural changes.

Layer	Technology	Notes
AI / Model	IBM watsonx.ai (preferred) OR Llama 3 / Mistral	Model-agnostic design. Swap without architectural change.
Embeddings	sentence-transformers (SBERT)	Open source. Local deployment preferred.
Clustering	UMAP + HDBSCAN + BERTopic	Open source pipeline.
Structured Storage	PostgreSQL	Signal, pattern, decision objects.
Raw Storage	S3 / Ceph / object storage	PDF ingestion, call transcripts.
Vector DB	Milvus (recommended) or Weaviate	Embedding storage and similarity search.
Processing	Python (primary), FastAPI (microservices)	
Orchestration	Airflow or Prefect	Pipeline scheduling and monitoring.
Agent Logic	IBM watsonx.ai (primary), LangChain (light use only)	LangChain for targeted tasks, not orchestration.
Ingestion Connectors	Airbyte	CRM, Slack, and structured source connectors.
Delivery	Slack API (Prism bot), email digest, Notion API or Confluent KB API	

V1 Implementation Plan (12–14 Weeks)

Phase	Weeks	Deliverables	Key Risk
0. Data Access	1–2	Map each signal source against access requirements. Identify IT/legal approvals. Obtain sample datasets. Decision: V1 vs. V2 source list.	Sources requiring >4 weeks of access negotiation are de-scoped to V2.
1. Scope & Config	2–4	Define 3 decision models. Build initial positioning assumption config with governance structure. Name decision owners. Define Decision Output format. Set up outcome tracking schema.	Decision owner designations must be confirmed before Phase 2 begins.
2. Signal Pipeline	4–7	Build ingestion for V1-approved sources (competitor scraping, analyst RSS, one internal source). Entity extraction, tagging pipeline, initial embeddings.	Internal source access (ISC, Slack) may slip. Plan for competitor + analyst only as minimum viable V1.
3. Quality + Interpretation	7–9	Source independence scoring. Recency weighting. Signal tier assignment. BERTopic clustering. LLM pattern compression with restricted prompts.	Manual review gate active for all pattern outputs.
4. Decision Engine	9–11	Rule-based decision engine for 3 models. Decision Output format implementation. Constraint config integration. Outcome tracking interface.	Positioning config must be reviewed and validated before decision logic is enabled.
5. Delivery + Calibration	11–14	Prism Slack bot. Weekly digest. 2-week live run with PMM leaders. Threshold calibration. Documentation.	Calibration requires active PMM participation. Schedule upfront.

V2 Capabilities (Post V1 Calibration)

- Predictive signal modeling: anticipate pattern shifts before threshold is crossed
- Narrative simulation: model downstream impact of positioning changes
- Automated outcome correlation: link decision choices to pipeline and win rate data
- Drift detection: identify when the positioning assumption config is diverging from market reality

Risk Register

This register captures the risks most likely to cause system failure or organizational non-adoption. Acknowledging these risks is not a case against building; it is the case for building with the right structure. Each risk has a defined mitigation.

Risk	Likelihood	Impact	Mitigation
Data access takes longer than expected	HIGH	HIGH	Phase 0 exists specifically to surface this. Sources requiring >4 weeks of access negotiation are V2 by default.
PMM treats Prism outputs as optional reading	HIGH	CRITICAL	Organizational adoption section. Leadership compact and decision SLA must be agreed before build begins. Without this, build should not proceed.
Positioning config drifts and becomes stale	MEDIUM	HIGH	Version control, change log, quarterly review cadence, and STALE flag built into Layer 5 governance.
Clustering surfaces spurious or false patterns	MEDIUM	HIGH	Manual review gate for first 30 days. Models with >40% Inaccurate outcome reviews in 90-day window are suspended.
LLM hallucination in interpretation layer	MEDIUM	MEDIUM	Restricted output format (X→Y only). Human spot-check during calibration phase. Format rejection rule rejects non-compliant outputs.
Same-source signals counted as independent	MEDIUM	MEDIUM	Source independence scoring in Layer 2. Same-origin signals count as one regardless of volume. Adversarial discount on competitor content.
IBM watsonx.ai model quality insufficient for pattern tasks	LOW	MEDIUM	Architecture is model-agnostic. Swap does not require structural changes. Evaluate during calibration phase.
Decision SLA not honored in practice	MEDIUM	HIGH	DEFERRED status is visible in leadership reporting. Governance compact makes non-response a logged, visible action rather than silence.

User Experience (Mockup with Fusion as product example)

Prism PMM decision intelligence

Week of May 26 | Integrations ↗

Open decisions	Deferred	Decided this quarter	Avg. decision time
3	0	4	3.2d

Competitive axis shift
Fusion AI positioning
ACTION REQUIRED
CONFIRMED
Commit by June 7

Entry point validity
Fusion pipeline entry drivers
MONITORING EMERGING
Commit by July 3

Narrative drift detection
Fusion platform vs. product framing
ACTION REQUIRED
ESTABLISHED
Commit by June 21

COMPETITIVE AXIS SHIFT

Fusion AI positioning

ACTION REQUIRED HIGH

PATTERN DETECTED

Market shifting from AI experimentation to AI operationalization. Enterprise buyers are evaluating platforms on production readiness, not experimentation capability.

CONSTRAINT

Fusion narrative assumes early-stage AI adoption framing. This assumption no longer holds.

SIGNAL BASIS

4	14d	3
signals	window	categories

CONFIRMED

Analyst Competitor Field

DECISION TYPE	COMMIT BY	DECISION OWNER
MESSAGING ACCENT SHIFT	June 7	PMM Lead, Platform Narrative

THE DECISION

A Commit and act

Shift Fusion messaging to emphasize operational AI readiness. Maintain platform narrative; update proof points and talk tracks to reflect production-stage AI workloads.

Choose A

B Hold — active choice

Hold current exploratory AI framing.

↳ RE-EVALUATE IF:
More than 2 field calls flag a too-early-stage objection within 30 days, OR any Tier 1 analyst publishes operational AI maturity as a primary buying criterion.

Choose B (hold with trigger)

IMPACT IF NO DECISION BY JUNE 7

Area: AI, platform, and sovereignty messaging

Timeframe: Q3 (medium)

Fusion positioned as exploratory in a market that has moved to operational accountability framing — visible to enterprise architects before Q3 campaigns launch.

Status: OPEN — awaiting decision from PMM Lead, Platform Narrative

Mark deferred

Add decision model ↗

Data Schema

Signal Object

```
{
  "id":               "string",
  "source":           "competitor | analyst | sales | field | slack | internal",
  "source_family":    "string",           // for independence deduplication
  "timestamp":        "datetime",
  "content":          "string",
  "entities":         ["product", "company", "theme"],
  "tags":             ["AI", "pricing", "governance", "hybrid"],
  "independence_wt": 0.5 | 0.75 | 1.0,  // assigned by Layer 2
  "recency_wt":       0.0 | 0.5 | 0.75 | 1.0,
  "weighted_score":  0.0                 // independence_wt * recency_wt
}
```

Pattern Object

```
{
  "id":           "string",
  "signal_ids":   ["signal_id", ...],
  "summary":      "Market is shifting from [X] to [Y] because [reason]",
  "direction":    "X_to_Y",
  "signal_tier":  "EMERGING | CONFIRMED | ESTABLISHED",
  "source_count": 3,                       // weighted unique source categories
  "created_at":   "datetime",
  "reviewed":     false                    // manual review gate flag
}
```

Constraint Object

```
{
  "id":                "string",
  "pattern_id":        "pattern_id",
  "assumption_id":     "config_assumption_id",
  "assumption_text":   "string",
  "broken_because":    "string",
  "affected_area":     "narrative | product | pricing | field"
}
```

Decision Object

```
{
  "id":               "string",
  "model":            "narrative_drift | entry_validity | axis_shift",
  "signal":           "string",
  "constraint":       "string",
  "decision_type":    "FULL_REPOSITION | MESSAGING_ACCENT | FRAMING_SHIFT
                       | HOLD_WITH_TRIGGER",
  "option_a":         "string",
  "option_b":         "string",
  "b_trigger":        "string | null",    // required if decision_type = HOLD
  "impact":           "string",
  "urgency":          "low | medium | high",
  "time_horizon":     "short | medium | long",
  "commit_by":        "date",
  "decision_owner":   "string",
  "status":           "OPEN | DECIDED-A | DECIDED-B | DEFERRED | EXPIRED",
  "decided_at":       "datetime | null",
  "notes":            "string | null",
  "review_30d":       "accurate | inaccurate | inconclusive | null",
  "review_60d":       "accurate | inaccurate | inconclusive | null",
  "review_90d":       "accurate | inaccurate | inconclusive | null"
}
```