

PRISM

The PMM Decision Intelligence Agent

Signals aren't the problem — PMM has more of them than any previous era. The problem is decision latency: the gap between when a shift is visible and when PMM commits to a response.

WHY DECISIONS STALL TODAY



Fragmented across 10+ systems

ISC, EPM, MCC, Slack, call transcripts, analyst papers, competitor content — no unified interpretation layer.



Manual, inconsistent detection

Pattern-spotting depends on which analyst happens to be paying attention that week.



“Seen” and “decided” look identical

Nothing distinguishes a signal PMM reviewed and held from one it never processed.



Deferral is free — until it isn't

The cost of not deciding is invisible right up until the market window has already closed.

*The consequence is not a lack of information. **It's decision latency — and that's where competitors gain ground.***

A decision-forcing engine — not a dashboard

Six layers turn scattered market signal into one forced, timed, owned decision.



SIGNAL QUALITY GATES EVERY LAYER — SOURCE INDEPENDENCE, RECENCY DECAY, AND ADVERSARIAL DISCOUNTING PREVENT PATTERN CONTAMINATION

EVERY DECISION TRIGGER INCLUDES

- ✓ An auto-calculated commit-by date
- ✓ A named decision owner — a role, never “PMM”
- ✓ Two concrete options: act (A), or hold with a defined trigger (B)
- ✓ An outcome log: OPEN → DECIDED → DEFERRED → reviewed at 30/60/90 days

Hold is not neutral. An unresolved trigger after commit-by is logged DEFERRED — visible to leadership.

THREE V1 DECISION MODELS



Narrative Drift

Is the field quietly repositioning the platform as a product instead of a platform?



Entry Point Validity

Does the emphasized entry point still match what’s actually driving new pipeline?



Competitive Axis Shift

Has the basis of competition moved — e.g. infrastructure vs. AI-system readiness?

Working Prototype Located At: <https://prism.stanilovic.net/>

Built to prove its own value

Outcome tracking and calibration are V1 requirements, not V2 enhancements — so we know if Prism is right.



3.2 days

avg. time to a logged decision once triggered



3

decision models live at V1 launch



12–14 wks

to the first production decision trigger



90 days

review cycle — >40% inaccurate suspends a model

PROOF POINT — FUSION AI POSITIONING

Signal: reaches CONFIRMED tier in 14 days across 3 source categories (analyst, field, competitor).

Decision forced: 7 days before the Q3 campaign brief locks — not after.

Impact avoided: Fusion stays framed as “exploratory AI” while the market has already moved to operational-readiness criteria.

WHAT WE NEED FROM LEADERSHIP TO BUILD IT



A named decision owner for each of the 3 V1 models



A decision SLA: 48-hr acknowledgment, 5 days to log status



“No action is a choice” — DEFERRED is visible to leadership



A signed governance compact: silence is not an accepted response

The architecture is scoped. The risk is mapped. Approve Prism to ship the first forced decision in 12 weeks.